RAG Assignment: Baseline Retrieval vs. Cross-Encoder Re-Ranking

Course: Large Language Models & Multimodal Applications (LMA)
Deliverable: Technical Report (Concise 4-Page Submission)
Code Repository: https://github.com/adityaND19/LMA-Assignment-1
Target Corpus: Regulation (EU) 2016/679 (GDPR) — Official Journal L 119 (88 pages, 444 sentence-aligned chunks)
Models: Bi-Encoder: `sentence-transformers/all-MiniLM-L6-v2` | Cross-Encoder: `cross-encoder/ms-marco-MiniLM-L-6-v2` | Generator: Google Gemini `gemini-3.1-flash-lite`

1. Architecture Comparison: Bi-Encoder vs. Cross-Encoder

The core architectural divergence between a Bi-Encoder (Pipeline 1) and a Cross-Encoder (Pipeline 2) lies in **how and when token interactions are computed**.

```
BI-ENCODER (Dual-Tower, Independent Embeddings):
  Query q   -----> [ Transformer Enc ] -----> u = Pool(H_q) ----\
                                                                Cosine Sim: (u . v) / (|u| |v|)
  Passage d -----> [ Transformer Enc ] -----> v = Pool(H_d) ----/
  * Offline: All passages d_i pre-embedded into FAISS index in O(1) query time. Zero cross-attention.

CROSS-ENCODER (Single-Tower, Joint Self-Attention):
  [CLS] + Query q + [SEP] + Passage d -----> [ Full Joint Transformer Self-Attention ] -----> Score s(q,d)
  * Online: Every token in q directly attends to every token in d across all 6 layers simultaneously.
```

1.1 Bi-Encoder (Independent Dual-Tower Representation)

- **Mechanism:** The query q and document chunk d are encoded independently: $$\mathbf{u} = \text{MeanPool}(\text{Enc}(q)), \quad \mathbf{v} = \text{MeanPool}(\text{Enc}(d))$$ Relevance is computed via cosine similarity: $\text{sim}(q, d) = (\mathbf{u} \cdot \mathbf{v}) / (|\mathbf{u}|_2 |\mathbf{v}|_2)$.
- **Advantages:** Highly scalable. Chunk vectors $\mathbf{v}_i$ are pre-computed offline and indexed into sub-millisecond Approximate Nearest Neighbor (ANN) indexes like FAISS `IndexFlatIP`.
- **Bottlenecks:** Compressing 100–300 words into a single 384-dimensional vector creates an information bottleneck. The query cannot influence how passage tokens are weighted. Consequently, dense bi-encoders struggle with:
- *Lexical mismatch and terminology gaps* (e.g., query "subcontractor" vs. statutory text "another processor").
- *Negations and qualifiers* (e.g., "SHALL NOT engage" vectors map adjacent to "SHALL engage" because the lexical context is nearly identical).
- *Distractor recitals* that repeat high-level keywords without containing operative rules.

1.2 Cross-Encoder (Joint Token Self-Attention)

- **Mechanism:** The query and passage are concatenated into a single input sequence: $$\mathbf{x} = [\text{CLS}], q_1, \dots, q_m, [\text{SEP}], d_1, \dots, d_n, [\text{SEP}]$$ The sequence passes through all transformer layers together. Every query token computes scaled dot-product self-attention with every passage token: $$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right)\mathbf{V}$$ The $[\text{CLS}]$ token is linearly projected to output a scalar relevance score $s(q, d) \in \mathbb{R}$.
- **Advantages:** Captures fine-grained word-level dependencies, modifier bindings, and conditional clauses. It easily recognizes whether a prohibition binds a "processor" or a "controller", or if an action requires "prior written authorisation".

2. Qualitative & Rank Shift Analysis

2.1 Failure Mode Categories & Rank Shift Summary

We evaluated 16 targeted GDPR stress queries across 5 distinct failure categories. The candidate pool consists of the Top-10 dense chunks retrieved in Stage 1.

Query ID	Category	Query Text	Critical Chunk (Statutory Basis)	Dense Rank	CE Rank	Retrieval Outcome
Q01	Keyword Trap	*Can companies share user passwords with third parties?*	Article 4 (`C187`)	#5	**#1**	**RE-RANK RESCUE**
Q02	Role Limit	*Can a data processor hire another subcontractor without permission?*	Article 28(2) (`C259`)	#4	**#1**	**RE-RANK RESCUE**
Q03	Exception	*Is processing sensitive data ever allowed without explicit consent?*	Article 9 (`C207` – `C210`)	>#10	—	Both miss (ceiling)
Q04	Exception	*Can a controller transfer personal data outside the EU?*	Article 44 (`C311`)	#5	#5	Both miss (in Top 10)
Q05	Negation	*What is explicitly prohibited when processing children data online?*	Article 8 (`C205`)	#5	#5	Both miss (in Top 10)
Q06	Negation	*When is automated decision-making NOT allowed?*	Article 22 (`C246`)	#1	#2	Both hit
Q07	Buried Chunk	*What is data protection by design and by default?*	Article 25 (`C253`)	#5	**#1**	**RE-RANK RESCUE**
Q08	Role Limit	*Who must appoint a Data Protection Officer?*	Article 37 (`C287`)	#1	#1	Both hit
Q09	Buried Chunk	*What information must be provided when a person requests access?*	Article 15 (`C232`)	#4	**#3**	**RE-RANK RESCUE**
Q10	Buried Chunk	*What penalties exist for serious GDPR breaches?*	Article 83 (`C418` , `C420`)	#2	#1	Both hit
Q11	Keyword Trap	*How long can companies keep customer data?*	Article 5 (`C196`)	#3	#3	Both hit
Q12	Exception	*Can personal data be processed for scientific research without consent?*	Article 89 (`C429`)	>#10	—	Both miss (ceiling)
Q13	Negation	*What rights can be restricted under national law?*	Article 23 (`C248`)	#1	#1	Both hit
Q14	Role Limit	*What breach notification obligations apply to processors?*	Article 33 (`C273`)	#1	#2	Both hit
Q15	Buried Chunk	*Under what conditions can an authority impose a temporary ban?*	Article 58 (`C354`)	#1	#2	Both hit
Q16	Keyword Trap	*Do users have the right to delete their account data?*	Article 17 (`C236`)	#3	#2	Both hit

Key Metric: Re-ranking increases Top-3 relevance hit rate from **43.8% (7/16)** to **68.8% (11/16)**, yielding a **+25.0% absolute improvement**.

2.2 Deep-Dive: Failure Queries and Answer 1 vs. Answer 2 Comparison

Failure Query 1: Q02 — Role Limits & Subcontracting

- **Query:** *"Can a data processor hire another subcontractor without permission?"*
- **Rank Shift:** `C259` (Article 28(2) Processor Obligations): **Dense Rank #4 (score: 0.531) $\rightarrow$ Cross-Encoder Rank #1 (score: +0.59)**
- *Baseline Top 3:* `C265` (Recital/Security), `C127` (Recital/Cooperation), `C033` (Recital/Contracts). All 3 chunks were general non-operative recitals.
- *Re-ranked Top 3:* `C259` (Article 28(2)), `C114` (Recital 114), `C262` (Article 28(4) Sub-processor liability).
- **Comparison of LLM Generations:**
- **Answer 1 (Baseline):** *"The provided context does not answer this question."*
  - *Analysis:* **Complete Omission / Refusal.** Because the operative rule sat at rank #4, the baseline prompt contained only background recitals. Strict adherence to prompt constraints forced the model to state the context lacked the answer.
- **Answer 2 (Re-ranked):** *"The processor SHALL NOT engage another processor without prior specific or general written authorisation of the controller [C259]. In the case of general written authorisation, the processor SHALL inform the controller of any intended changes concerning the addition or replacement of other processors, thereby giving the controller the opportunity to object to such changes [C259]."*
  - *Analysis:* **Legally Accurate.** Cross-Encoder promoted `C259` from #4 to #1, enabling the LLM to directly cite the statutory prohibition and conditional exception.

Failure Query 2: Q07 — Buried Chunk & Core Regulatory Definitions

- **Query:** *"What is data protection by design and by default?"*

- **Rank Shift:** `C253` (Article 25 Data Protection by Design and Default): **Dense Rank #5 (score: 0.589) $\rightarrow$ Cross-Encoder Rank #1 (score: +8.13)**
- *Baseline Top 3:* `C085` (Recital 85), `C091` (Recital 91), `C323` (Binding Corporate Rules). Flooded with high-level recitals mentioning "principles".
- *Re-ranked Top 3:* `C253` (Article 25), `C085` (Recital 85), `C084` (Recital 84 on organizational measures).
- **Comparison of LLM Generations:**
- **Answer 1 (Baseline):** *"The provided context does not answer this question."*
  - *Analysis:* **Information Miss.** The dense retrieval matched recitals mentioning data protection principles generally, failing to deliver the operative definition.
- **Answer 2 (Re-ranked):** *"Data protection by design and by default requires the controller to implement appropriate technical and organisational measures, such as pseudonymisation, to effectively integrate data-protection principles like data minimisation into the processing [C253]. This involves ensuring that, by default, only personal data necessary for each specific purpose of the processing are processed [C253]..."*
  - *Analysis:* **Complete & Precise.** Cross-Encoder promoted `C253` from #5 to #1, capturing both paragraphs of Article 25.

Failure Query 3: Q09 — Buried Chunk & Disclosure Checklist

- **Query:** *"What information must be provided when a person requests access to their data?"*
- **Rank Shift:** `C232` (Article 15 Right of Access): **Dense Rank #4 (score: 0.598) $\rightarrow$ Cross-Encoder Rank #3 (score: +2.48)**
- **Comparison of LLM Generations:**
- **Answer 1 (Baseline):** Stated the context did not specify the required content, only citing ancillary Recital 66 (`C066`) regarding excessive quantities of data.
- **Answer 2 (Re-ranked):** Accurately enumerated the full statutory checklist under Article 15 (`C232`): purposes of processing, categories of data, recipients, envisaged storage period, and rights to rectification/erasure.

2.3 The Retrieval Ceiling Limitation

In **Q03** (Article 9 sensitive data exceptions) and **Q12** (Article 89 research derogations), the critical chunks were ranked at dense positions >#10 (e.g. `C208` was dense #19). Because Stage 1 only surfaces the Top 10 candidates, the Cross-Encoder never observed these passages. **Finding:** *A re-ranker cannot rescue chunks that fall outside the first-stage candidate pool. The recall of Stage 1 forms the absolute upper bound of Stage 2.*

---

3. Latency Profiling

Retrieval and generation latencies were measured on CPU with pre-warmed models across $N=320$ retrieval runs and $N=16$ LLM calls using monotonic timers (`time.perf_counter()`).

Pipeline / Execution Stage	Median (ms)	Mean (ms)	p95 (ms)	Samples (n)	Overhead vs. Dense
Pipeline 1: Dense Retrieval (Query encode + FAISS Top-10 + Top-3 slice)	**20.4**	17.1	21.6	320	1.0x (Baseline)
Re-Ranking Step Only (Cross-Encoder on 10 pairs + sort)	**424.0**	391.8	437.8	320	+424 ms
Pipeline 2: Full Two-Stage Retrieval (Dense Top-10 + Re-Ranking)	**443.8**	408.8	459.0	320	21.7x
LLM Generation: Answer 1 (Baseline Top-3 $\rightarrow$ Gemini API)	**1,351.1**	2,412.7	6,807.1	16	—
LLM Generation: Answer 2 (Re-ranked Top-3 $\rightarrow$ Gemini API)	**1,330.0**	2,269.0	7,391.3	16	—

Latency Trade-Off Analysis

1. **Retrieval vs. Generation:** While re-ranking adds ~424 ms of CPU time, end-to-end response time is dominated by LLM token generation (~1.3–2.4 s).
2. **Interactive Feasibility:** Adding 424 ms increases total round-trip time from ~1.37 s to ~1.77 s (~29% increase), well within standard user tolerances (sub-2 seconds) while delivering a **+25.0% boost in answer accuracy**.

---

4. Scaling Question: Why Not Run a Cross-Encoder Over the Whole Collection?

A central system design question is: *Why not run a Cross-Encoder over the entire document corpus from the very start, eliminating dense retrieval entirely?*

4.1 Asymptotic Complexity Comparison

Let N be corpus chunks, L_q query length, L_d chunk length, and d embedding dimension.

1. **Bi-Encoder Scaling:**
2. *Offline (once):* Chunks are encoded once into vector space: $O(N \cdot L_d^2)$.
3. *Online (query time):* Query is encoded once: $O(L_q^2)$. Nearest-neighbor search over pre-computed vectors requires $O(N \cdot d)$ dot products, reducible via ANN indexing (FAISS HNSW/IVF) to sub-linear time: $$\text{Query Complexity}_{\text{Bi-Encoder}} = O(L_q^2 + \log N \cdot d)$$
4. **Exhaustive Cross-Encoder Scaling:**
5. Query and document are concatenated: $(L_q + L_d)$ tokens.
6. Attention across all layers has quadratic token complexity $O((L_q + L_d)^2)$.
7. Because cross-attention dynamically couples query and chunk tokens, **no pre-computation or indexing is possible**. Every query requires evaluating all N chunks at runtime: $$\text{Query Complexity}_{\text{Exhaustive CE}} = O(N \cdot (L_q + L_d)^2)$$

4.2 Empirical Scaling Benchmark

Using our measured CPU throughput of ~41.0 ms per (query, chunk) pair:

Collection Scale	Chunk Count (N)	Exhaustive Cross-Encoder Latency	Two-Stage Funnel (Dense + Top-10 CE)	Speedup Factor
This Experiment (GDPR)	444	**18.2 seconds**	**0.44 seconds**	**41.4x**
Enterprise Knowledge Base	10,000	**410.3 seconds (~6.8 min)**	**~0.45 seconds**	**~910x**
Web / Legal Corpus	1,000,000	**41,032.8 seconds (~11.4 hours)**	**~0.46 seconds**	**~89,000x**

4.3 Conclusion: The Two-Stage Funnel as Optimal Architecture

Exhaustive cross-encoder retrieval is computationally impossible for interactive RAG systems. The **Two-Stage Funnel** provides the optimal Pareto frontier: - **Stage 1 (Bi-Encoder):** Filters $N = 10^6 \rightarrow K = 10$ in $O(\log N)$ time (<25 ms). - **Stage 2 (Cross-Encoder):** Scores only $K = 10$ candidate pairs with constant $O(K)$ latency (~424 ms).

This guarantees the semantic precision of joint self-attention at constant, scalable sub-second latency.